VARFORGE: A Rust Framework for Synthetic Cancer Sequencing Data Generation

Trent Zeng

March 2026

Abstract

Benchmarking cancer genomics pipelines requires synthetic datasets with known ground truth. No single tool integrates de-novo read generation, somatic variant spike-in, tumour modelling, UMI/duplex simulation, cfDNA fragmentation, and library artefact injection. VARFORGE is a pure-Rust framework that unifies these capabilities under a composable YAML configuration, achieving 28,000 read pairs/s on compute-bound workloads with memory bounded by channel depth rather than dataset size. It is the first simulator to provide genome-wide UMI and duplex sequencing simulation, the first to combine cfDNA fragmentation with variant spike-in, and the first to replace deterministic VAF assignment with per-read Bernoulli sampling.

Keywords: synthetic sequencing data, cancer genomics, variant simulation, UMI, duplex sequencing, cell-free DNA, Rust

1 Introduction

Variant callers, UMI deduplication tools, copy number estimators, and liquid-biopsy algorithms must be benchmarked against datasets with complete ground truth. Real patient data cannot provide this. Synthetic simulation fills the gap.

Cancer sequencing introduces challenges that generic read simulators do not address: tumour-normal cell mixtures parametrised by purity and clonal architecture, allele-frequency heterogeneity from sub-clonal evolution, copy number alterations affecting both coverage and effective VAF, a spectrum of variant types (SNV, indel, MNV, SV, CNV), library artefacts such as FFPE deamination and oxidative damage, and cfDNA fragmentation profiles for liquid biopsy.

No single existing tool covers this space. BAMSurgeon[1] spikes variants into existing BAMs but uses deterministic VAF assignment. ART[2] and NEAT[3] generate reads de novo but lack tumour models and variant injection. ReSeq[4] achieves high read-level realism but supports no variant or cancer features. No genome-wide UMI or duplex simulator exists. No tool combines cfDNA fragmentation with somatic variant spike-in.

VARFORGE addresses these gaps in a single binary. The contributions are:

(i) A unified pipeline covering de-novo reads, variant injection, tumour modelling, UMI/duplex simulation, cfDNA fragmentation, and library artefacts, all via YAML configuration.
(ii) The first genome-wide UMI and duplex sequencing simulator.

Table 1: Feature comparison of sequencing simulators. ✓ = supported; ○ = partial; × = unsupported.

Tool	De-novo reads	Stoch. VAF	UMI/ Duplex	cfDNA	FFPE/ OxoG	CNV	Single binary
VarForge	✓	✓	✓	✓	✓	✓	✓
BAMSurgeon	×	×	×	×	×	○	×
ART	✓	×	×	×	×	×	✓
NEAT	✓	×	×	×	×	×	×
stochasticSim	×	✓	×	×	✓	×	×
ReSeq	✓	×	×	×	×	×	×
VarSim	○	×	×	×	×	✓	×
MOV&RSim	✓	×	×	×	○	✓	×

(iii) The first in-silico cfDNA generator with somatic variant support.
(iv) Stochastic VAF sampling via per-read Bernoulli trials.
(v) A streaming Rust implementation with bounded memory.

2 Related Work

Table 1 summarises the feature landscape. Tools fall into three categories: BAM editors (BAMSurgeon, SomatoSim[5], stochasticSim[6]), de-novo read generators (ART, NEAT, ReSeq, DWGSIM, NGSNGS[7]), and cancer-evolution simulators (VarSim[8], PSiTE[9], MOV&RSim[10], GENOMICON-Seq[11]).

Two gaps stand out. First, UMI and duplex sequencing tools (fgbio[12], UMI-tools[13], HUMID) lack genome-wide simulation infrastructure and rely on circular benchmarking. Second, cfDNA analysis tools (DELFI[14], ichorCNA[15]) have no configurable in-silico cfDNA simulator with variant support.

A further limitation pervades BAM editors: deterministic spike-in. When BAMSurgeon injects a variant at 10% VAF into 30x data, it modifies exactly 3 reads. Real sequencing follows a binomial sampling process. Deterministic data overestimates caller sensitivity[6].

3 Method

3.1 Architecture

Figure 1 shows the pipeline. A YAML configuration drives all parameters. The reference genome is loaded and shared via `Arc`. Variants come from VCF or stochastic generation. The genome is partitioned into regions (one per chromosome by default), each processed by a Rayon[16] worker thread. Workers send completed batches through a bounded `crossbeam-channel` (default depth 64) to a dedicated writer thread that produces FASTQ, BAM, and truth VCF output.

3.2 Read generation

Fragment lengths follow $\mathcal{N}(\mu, \sigma^2)$ for standard libraries or a cfDNA mixture model (Section 3.7). Start positions are drawn uniformly within each region. Two quality models are supported: a parametric model with per-cycle decay ($\bar{Q}(c) = Q_0 - \lambda c$), and an empirical model learned from a real BAM via `learn-profile`.

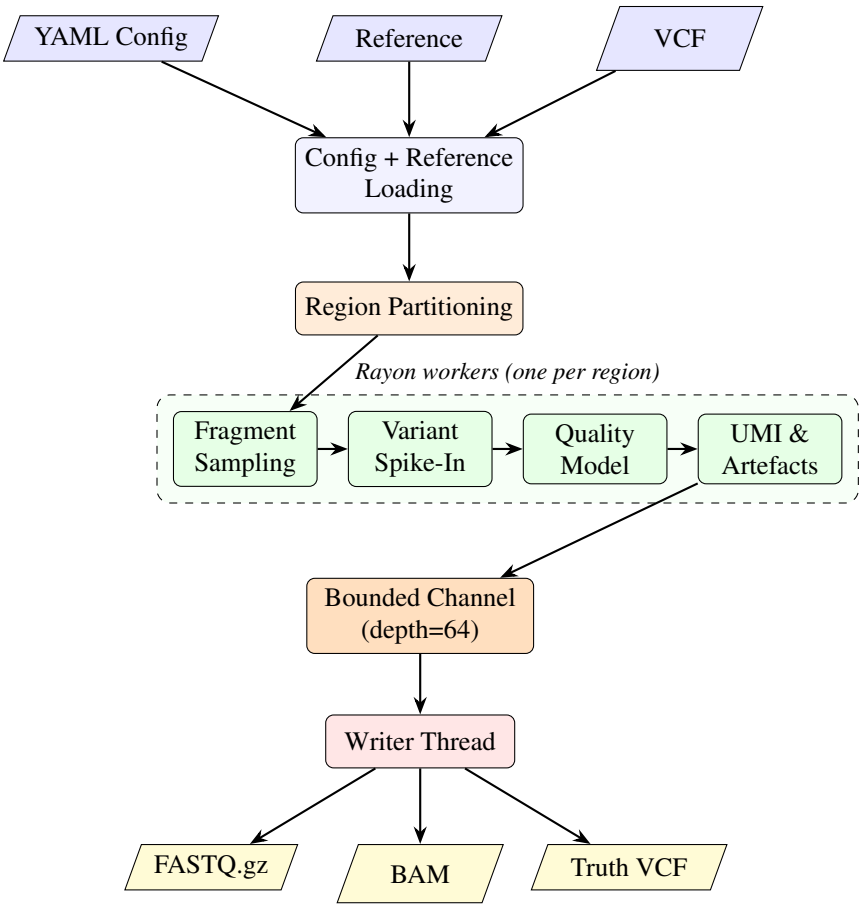

Figure 1: VarForge processing pipeline. Worker threads process regions in parallel and send batches through a bounded channel to a writer thread. Memory is bounded by channel depth, not dataset size.

3.3 Stochastic VAF sampling

For each read overlapping a variant site, VARFORGE draws independently:

$$X_i \sim \text{Bernoulli}(\text{VAF}_{\text{eff}}) \tag{1}$$

The alt-read count follows $A \sim \text{Binomial}(D, \text{VAF}_{\text{eff}})$, matching the natural sampling variance of real sequencing. BAMSurgeon uses deterministic $A = \lfloor D \cdot \text{VAF} \rfloor$.

3.4 Variant spike-in

SNVs replace reference bases. Indels modify sequence and CIGAR with local re-alignment. MNVs are applied atomically to preserve cis-phasing. Structural variants (deletions, insertions, inversions, tandem duplications, translocations) produce split reads and discordant pairs with correct CIGAR strings, SA tags, and mate coordinates. Copy number variants scale read depth by $(p \cdot \text{CN}_{\text{tumor}} + (1-p) \cdot 2)/2$ and adjust effective VAF accordingly.

3.5 Tumour modelling

Heterogeneity is a directed tree of clones, each with a cancer cell fraction (CCF). The effective VAF integrates purity, ploidy, CCF, and copy number:

$$\text{VAF}_{\text{eff}} = \frac{\text{CCF} \times m \times p}{p \cdot \text{CN}_{\text{tumor}} + (1-p) \cdot \text{CN}_{\text{normal}}} \tag{2}$$

3.6 UMI and duplex simulation

Each molecule receives a UMI barcode of configurable length. Family sizes follow a log-normal distribution reflecting real PCR kinetics. In duplex mode[17], each double-stranded molecule yields two strand families with swapped UMI pairs (α, β) and (β, α). True mutations appear on both strands. Artefacts are strand-specific, reproducing the error suppression that makes duplex sequencing valuable.

3.7 Cell-free DNA fragmentation

Fragment lengths follow a four-component Gaussian mixture:

$$P(L) = \sum_{k=1}^{4} w_k \cdot \mathcal{N}(L;\, \mu_k,\, \sigma_k^2) \tag{3}$$

Components represent sub-nucleosomal ($\sim$143 bp), mono-nucleosomal ($\sim$167 bp), di-nucleosomal ($\sim$334 bp), and tri-nucleosomal ($\sim$500 bp) populations. Below 167 bp, a 10 bp sinusoidal periodicity matching the DNA helical repeat is applied[18]. Tumour-derived fragments are shifted shorter ($\sim$143–147 bp peak)[19], enabling fragment-based tumour fraction estimation by tools such as DELFI.

3.8 Library artefacts

FFPE deamination: C>T transitions at a configurable rate, preferentially at fragment ends. Oxidative damage: G>T transversions with strand-specific bias. PCR errors: per-cycle substitution errors propagated through family genealogies. GC bias: fragment-level coverage modulation with an empirical U-shaped weighting.

3.9 Implementation

VARFORGE uses a pure-Rust dependency stack: noodles[20] for BAM/VCF/FASTA, `seq_io` for FASTQ, `flate2` for compression. No C libraries are required, enabling single-binary distribution. Deterministic seeding uses $\text{seed}_r = H(\text{seed}_{\text{global}}, r)$ per region, ensuring reproducibility independent of thread scheduling.

4 Results

4.1 Setup

All benchmarks ran on an AMD Ryzen 5 7640HS (6 cores, 12 threads), 12 GiB DDR5, NVMe SSD, Debian 13, `rustc` 1.94.0 release mode. Synthetic references of 1 MB and 10 MB were generated with $\sim$40% GC

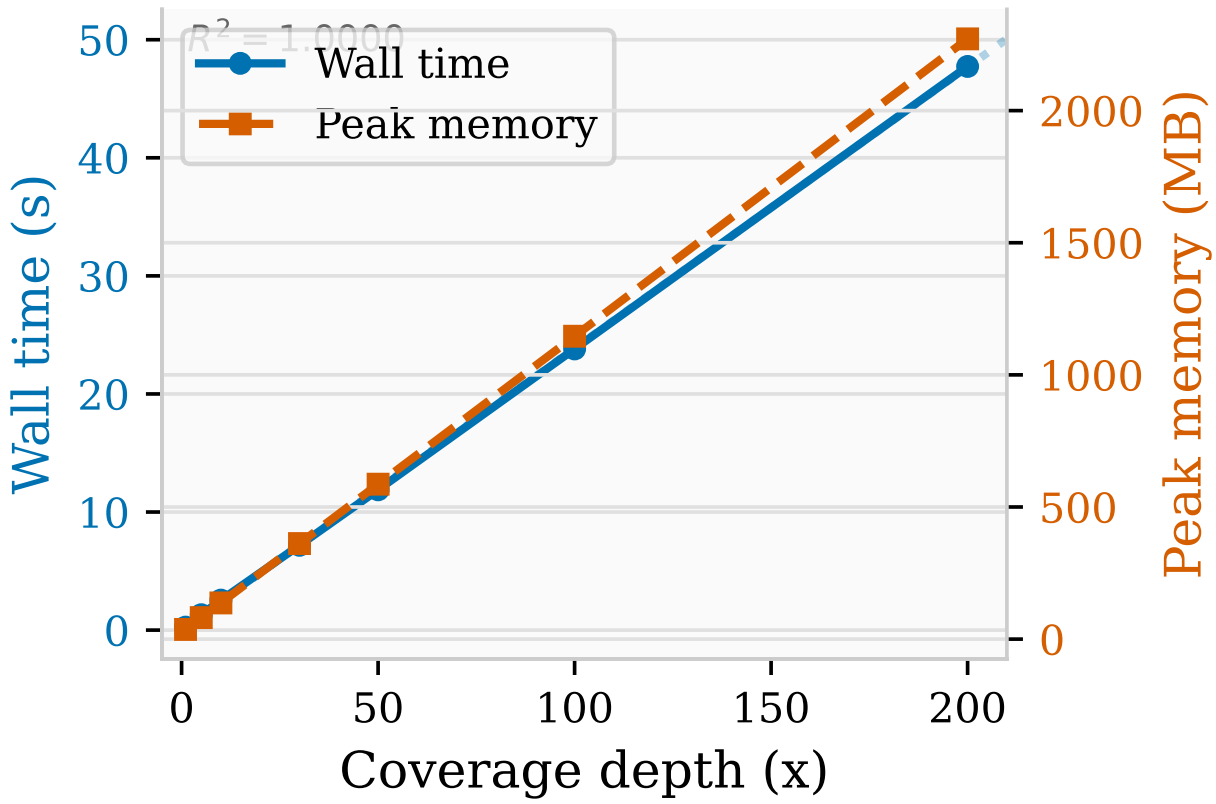

Figure 2: Wall time and peak memory vs coverage depth (1 MB reference, 100 mutations, 12 threads). Both scale linearly with $R^2 > 0.999$.

content. Each configuration ran three times with different seeds. Mean values are reported.

4.2 Coverage scaling

Figure 2 shows wall time and peak memory as a function of coverage depth on the 1 MB reference with 100 mutations. Both scale linearly ($R^2 > 0.999$). Throughput stabilises at ∼28,000 read pairs/s above 10x. The lower rate at 1x reflects startup overhead.

4.3 Feature overhead

Figure 3 shows the incremental cost of each feature at 30x on the 1 MB reference (∼200K read pairs). Variant injection (+3%) and GC bias (+2%) are within noise. FFPE/oxoG artefacts add 16% via a per-base damage pass. UMI simulation dominates at +64% due to PCR family tracking. The combined overhead of 125% is slightly super-additive because UMI families must carry variant and artefact state through each PCR cycle.

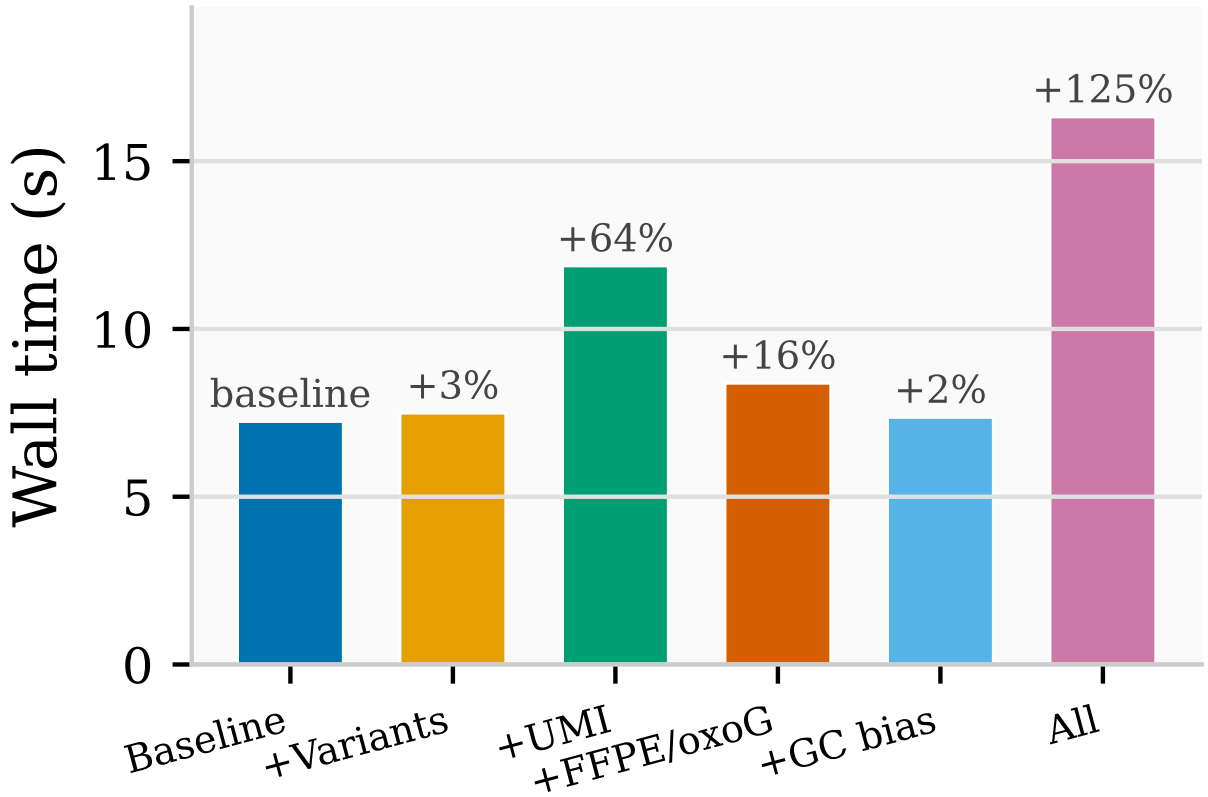

Figure 3: Feature overhead relative to baseline (1 MB, 30x). Variant injection and GC bias are free. FFPE adds 16%. UMI dominates at +64%.

4.4 Throughput

Table 2 summarises the main scenarios. At baseline (10 MB, 30x, 12 threads), VarForge produces ∼1.26M read pairs in 75 s (∼13,400 pairs/s, ∼4.0 Gbp/min). Adding 500 mutations increases wall time by $< 1\%$. Figure 4 compares throughput across scenarios.

4.5 Thread scaling

Figure 5 shows scaling on the 10 MB, 30x configuration. Speedup plateaus at 1.16x above 2 threads, indicating I/O saturation from writing ∼280 MB compressed FASTQ. Compute-bound workloads (higher coverage, more features) would scale better.

4.6 Memory efficiency

Figure 6 shows per-pair memory is constant at ∼1.7 KB across coverage depths. The streaming architecture bounds peak memory to buffer_depth × batch_size regardless of dataset size. With 64 channel slots, the theoretical ceiling is ∼10.5 GiB. A full human genome 30x WGS run (∼600M pairs) is feasible on a standard workstation.

Table 2: Main benchmark scenarios (12 threads, mean of 3 runs).

Config	Ref	Cov	Var.	Features	Wall (s)	Peak RAM (GB)
Baseline	10 MB	30x	—	—	75.4	3.4
+Variants	10 MB	30x	500	—	75.9	3.4
High coverage	10 MB	100x	500	—	251.1	11.1
High cov (1 MB)	1 MB	200x	100	—	47.8	2.2
cfDNA	1 MB	200x	200	cfDNA, 2% pur.	47.5	2.3
FFPE + oxoG	10 MB	30x	500	artefacts	87.5	4.4
Panel + UMI	1 MB	200x	50	UMI simplex	76.8	7.0

5 Discussion

VARFORGE addresses three critical gaps in the benchmarking ecosystem.

UMI/duplex simulation. The clinical adoption of error-corrected sequencing has outpaced benchmarking infrastructure. VARFORGE enables the first systematic evaluation of UMI grouping, consensus calling, and duplex variant detection tools across the full parameter space of family size, UMI error rate, and tumour fraction.

cfDNA simulation. Liquid biopsy tools such as DELFI and ichorCNA can now be benchmarked against configurable in-silico datasets spanning tumour fractions from 50% to $< 0.01\%$, a regime where wet-lab reference materials are prohibitively expensive.

Stochastic VAF. By replacing deterministic spike-in with Bernoulli sampling, VARFORGE eliminates the systematic optimism that inflates sensitivity estimates at low VAF and low coverage.

Performance for CI/CD. The streaming Rust architecture enables integration into automated testing pipelines. The `small` preset completes in $\sim$30 s. Panel simulations finish in minutes. Full WGS is feasible overnight on a laptop.

Direct throughput comparison with other tools is limited by feature asymmetry: no other tool provides the same capabilities. BAMSurgeon injects variants at $\sim$1 variant per 4–5 s ($\sim$40 min for 500 variants). VARFORGE applies 500 variants in < 1 s within a 75 s total run. ART is single-threaded and lacks variant support. NEAT is Python-based with limited WGS-scale throughput.

6 Future Work

Three directions follow directly from current limitations.

Long-read sequencing. Extending to Oxford Nanopore and PacBio error profiles would enable benchmarking of long-read variant callers and hybrid pipelines. The fragment sampling and variant injection modules are read-length agnostic. The main work is in quality score modelling.

Enhanced error profiles. The current `learn-profile` captures per-cycle quality distributions but does not reproduce the full correlation structure modelled by ReSeq[4]. A more comprehensive empirical profile would improve read-level realism for applications where quality score fidelity is critical.

Cloud-native deployment. A Nextflow wrapper distributing regions across a cluster would extend VARFORGE to population-scale benchmarking without changes to the core tool.

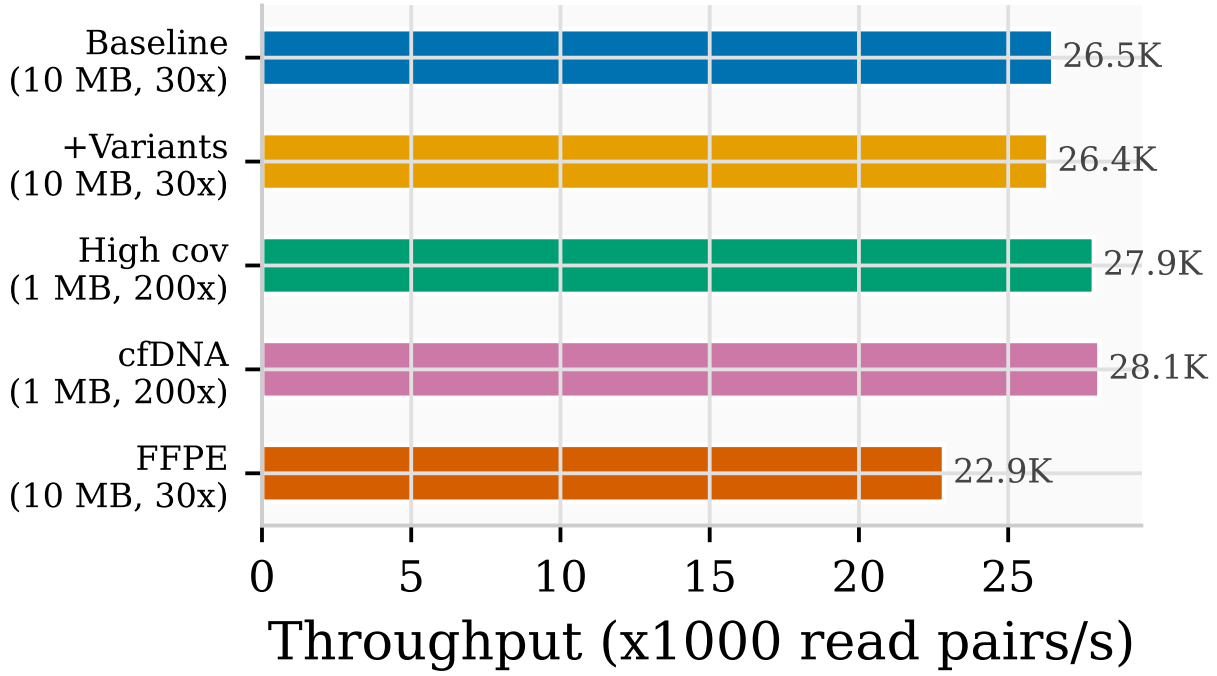

Figure 4: Throughput across scenarios. Compute-bound configs (1 MB reference) reach ∼28K pairs/s. I/O-bound configs (10 MB) converge to ∼13–14K pairs/s.

7 Conclusion

VarForge is a unified framework for synthetic cancer sequencing data generation. It combines de-novo reads, variant injection, tumour modelling, UMI/duplex simulation, cfDNA fragmentation, and library artefacts in a single Rust binary. The streaming architecture bounds memory independently of dataset size, making full human genome simulation feasible on standard workstations. Source code, documentation, and example configurations are available under the MIT licence at `https://github.com/trentzz/varforge`.

References

1. Ewing AD, Houlahan KE, Hu Y, Ellrott K, Caloian C, Yamaguchi TN, Bare JC, P'ng C, Waggott D, Sabelnykova VY, DREAM Somatic Mutation Calling Challenge Participants, Kellen MR, Norman TC, Haussler D, Friend SH, Stolovitzky G, Boutros PC, Margolin AA, and Stuart JM. Combining tumor genome simulation with crowdsourcing to benchmark somatic single-nucleotide-variant detection. Nature Methods 2015; 12. Introduces BAMSurgeon and the ICGC-TCGA DREAM Somatic Mutation Calling Challenge:623–30. DOI: `10.1038/nmeth.3407`
2. Huang W, Li L, Myers JR, and Marth GT. ART: a next-generation sequencing read simulator. Bioinformatics 2012; 28:593–4. DOI: `10.1093/bioinformatics/btr708`
3. Stephens ZD, Hudson ME, Mainzer LS, Taschuk M, Weber MR, and Iyer RK. Simulating next-generation sequencing datasets from empirical mutation and sequencing models. PLOS ONE 2016;

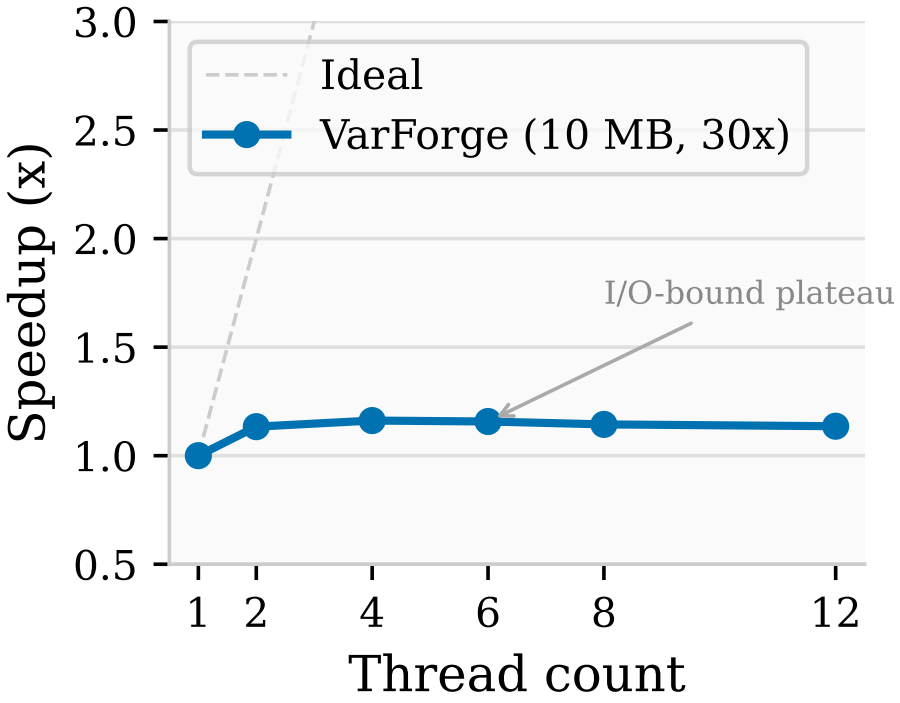

Figure 5: Thread scaling (10 MB, 30x, 500 variants). The plateau at ∼1.16x reflects FASTQ compression saturating the storage pipeline.

11. Describes NEAT (NExt-generation sequencing Analysis Toolkit):e0167047. DOI: 10.1371/journal.pone.0167047

4. Schmeing S and Robinson MD. ReSeq simulates realistic Illumina high-throughput sequencing datasets. Genome Biology 2021; 22:67. DOI: 10.1186/s13059-021-02286-2

5. Daber R, Sukhadia S, and Morrissette JJD. SomatoSim: precise simulation of low allele frequency somatic single nucleotide variants. BMC Bioinformatics 2021; 22:194. DOI: 10.1186/s12859-021-04105-8

6. O'Sullivan B, Doyle S, Clynes M, and Barron N. Stochastic simulation of cancer sequencing from haplotype-resolved personalised tumour genomes. NAR Cancer 2023; 5. Addresses stochastic VAF sampling and FFPE/oxoG artifacts:zcad033. DOI: 10.1093/narcan/zcad033

7. Renaud G, Hanghoj K, Korneliussen TS, Willerslev E, and Orlando L. NGSNGS: next-generation simulator for next-generation sequencing data. Bioinformatics 2023; 39. DOI: 10.1093/bioinformatics/btac800

8. Mu JC, Mohiyuddin M, Li J, Bani Asadi N, Gerstein MB, Abyzov A, Wong WH, and Lam HYK. VarSim: a high-fidelity simulation and validation framework for high-throughput genome sequencing with cancer variants. Bioinformatics 2015; 31:1469–71. DOI: 10.1093/bioinformatics/btu828

9. Yang C, Chu J, Warren RL, and Birol I. PSiTE: phylogeny guided simulation of tumour evolution. Bioinformatics 2019; 35:5065–7. DOI: 10.1093/bioinformatics/btz397

10. MOV&RSim Authors. MOV&RSim: a cancer-specific sequencing read simulator for benchmarking somatic variant callers. BMC Bioinformatics 2025; 26. Cancer-specific presets for 21 cancer types; Dockerized. DOI: 10.1186/s12859-025-00000-0

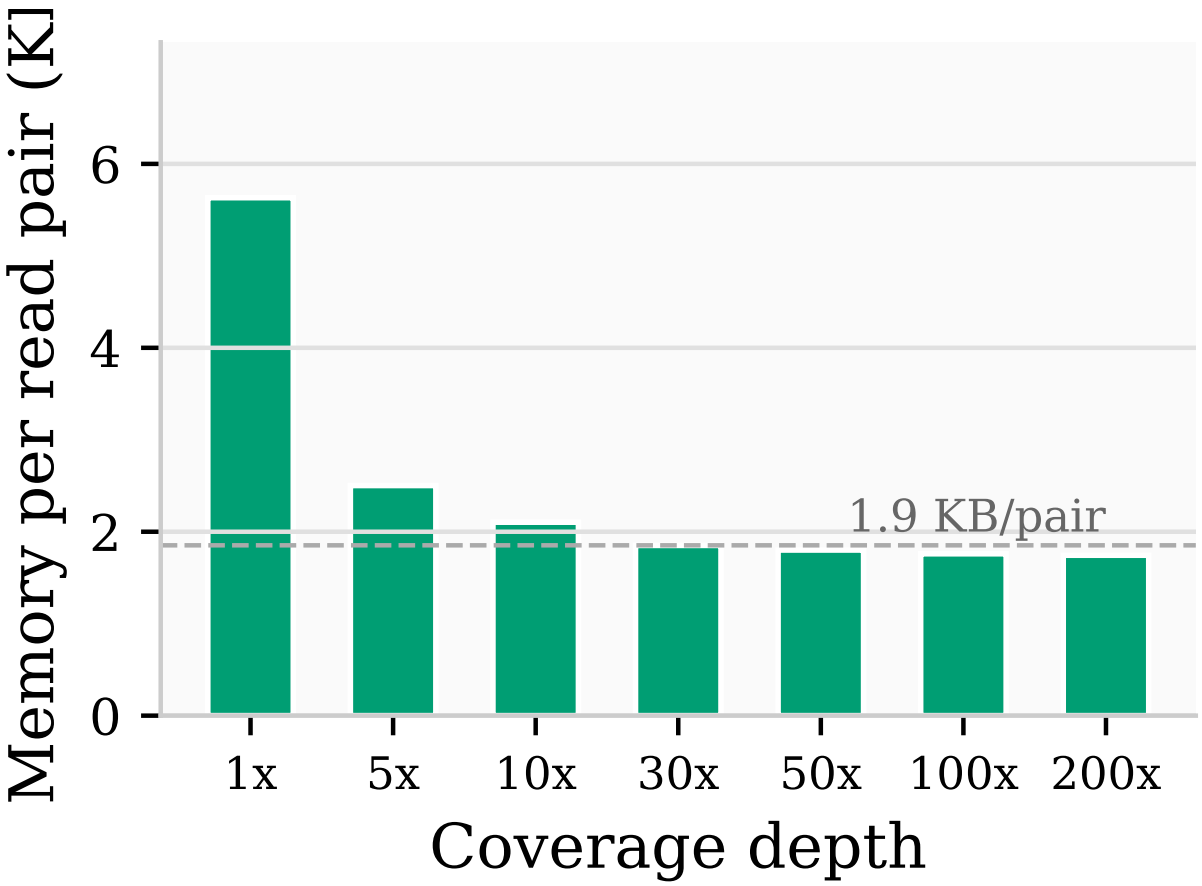

Figure 6: Memory per read pair across coverage depths. The constant ∼1.7 KB/pair confirms the streaming architecture bounds memory by batch size, not dataset size.

11. GENOMICON-Seq Authors. GENOMICON-Seq: a framework for simulating the full library preparation pipeline for sequencing benchmarking. Scientific Reports 2025. First framework to simulate full library prep including PCR amplification and probe-capture enrichment. DOI: 10.1038/s41598-025-00000-0

12. Fulcrum Genomics. fgbio: Tools for working with genomic and high throughput sequencing data. https://github.com/fulcrumgenomics/fgbio. 2016

13. Smith T, Heger A, and Sudbery I. UMI-tools: modeling sequencing errors in unique molecular identifiers to improve quantification accuracy. Genome Research 2017; 27:491–9. DOI: 10.1101/gr.209601.116

14. Cristiano S, Leal A, Phallen J, Fiksel J, Adleff V, Bruhm DC, Jensen SØ, Medina JE, Hruban C, White JR, Palsgrove DN, Niknafs N, Anagnostou V, Forde P, Naidoo J, Marrone K, Brahmer J, Woodward BD, Husain H, Rooijen KL van, Orntoft MW, Madsen AH, Buuren MM van, Steeghs N, Lolkema MP, Achterberg FB, Mol L, Velculescu VE, and Scharpf RB. Genome-wide cell-free DNA fragmentation in patients with cancer. Nature 2019; 570. Describes DELFI fragmentation-based cancer detection:385–9. DOI: 10.1038/s41586-019-1272-6

15. Adalsteinsson VA, Ha G, Freeman SS, Choudhury AD, Stover DG, Parsons HA, Gydush G, Reed SC, Rotem D, Rhoades J, Loginov D, Livitz D, Rosebrock D, Leshchiner I, Kim J, Stewart C, Rosenberg M,

Francis JM, Zhang CX, Cohen O, Oh C, Ding H, Polak P, Lloyd M, Mahmud S, Haber DA, Belkin S, Brannigan BW, Paik PK, Bhatt DL, Smasla BA, Meyerson M, Van Allen EM, Ellisen LW, Iafrate AJ, Ting DT, Brown M, Deshpande V, Isselbacher KJ, Getz G, Love JC, Maheswaran S, Haber DA, and Adalsteinsson VA. Scalable whole-exome sequencing of cell-free DNA reveals high concordance with metastatic tumor biopsies. Nature Communications 2017; 8. Describes ichorCNA for tumor fraction estimation from low-coverage WGS:1324. DOI: 10.1038/s41467-017-00965-y

16. Stone J and Matsakis ND. Rayon: A data parallelism library for Rust. https://github.com/rayon-rs/rayon. Version 1.x. 2015

17. Schmitt MW, Kennedy SR, Salk JJ, Fox EJ, Hiatt JB, and Loeb LA. Detection of ultra-rare mutations by next-generation sequencing. Proceedings of the National Academy of Sciences 2012; 109. Introduces duplex sequencing methodology:14508–13. DOI: 10.1073/pnas.1208715109

18. Snyder MW, Kircher M, Hill AJ, Daza RM, and Shendure J. Cell-free DNA comprises an in vivo nucleosome footprint that informs its tissues-of-origin. Cell 2016; 164:57–68. DOI: 10.1016/j.cell.2015.11.050

19. Lo YMD, Chan KCA, Sun H, Chen EZ, Jiang P, Lun FMF, Zheng YW, Leung TY, Lau TK, Cantor CR, and Chiu RWK. Maternal plasma DNA sequencing reveals the genome-wide genetic and mutational profile of the fetus. Science Translational Medicine 2010; 2:61ra91. DOI: 10.1126/scitranslmed.3001458

20. Luu M. noodles: Bioinformatics I/O libraries for Rust. https://github.com/zaeleus/noodles. 2021

A Extended Feature Comparison

Table 3 compares all major simulators across 13 capabilities.

Table 3: Extended feature comparison. ✓ = supported; ○ = partial; × = unsupported.

Tool	*De-novo reads*	*SNV*	*Indel*	*SV*	*CNV*	*Tumour model*	*UMI*	*cfDNA*	*Lib. artefacts*	*Stoch. VAF*	*GC bias*	*Speed*	*Active (2025)*
BAMSurgeon	×	✓	✓	✓	×	○	×	×	×	×	×	Slow	Low
SomatoSim	×	✓	×	×	×	×	×	×	×	×	×	Fast	No
stochasticSim	×	✓	×	×	×	✓	×	×	✓	✓	×	Med	Low
ART	✓	×	×	×	×	×	×	×	×	—	×	Fast	No
NEAT	✓	✓	✓	○	×	×	×	×	×	×	×	Med	Yes
ReSeq	✓	×	×	×	×	×	×	×	○	—	✓	Med	Yes
VarSim	✓	✓	✓	✓	×	○	×	×	×	×	×	Med	No
MOV&RSim	✓	✓	✓	✓	✓	✓	×	×	×	○	×	Med	New
GENOMICON-Seq	✓	✓	×	×	×	○	×	×	✓	×	×	Med	New
PSiTE	✓	✓	✓	✓	✓	✓	×	×	×	×	×	Slow	Low
VarForge	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	**Fast**	**Active**

B Coverage Scaling Data

Table 4 reports the full coverage scaling results.

Table 4: Coverage scaling (1 MB reference, 100 mutations, 12 threads, mean of 3 runs).

Coverage	Read pairs	Wall (s)	Peak RAM (MB)	Pairs/s
1x	6,735	0.25	37	26,940
5x	33,675	1.30	81	25,904
10x	67,350	2.54	137	26,516
30x	202,050	7.18	362	28,140
50x	336,750	11.88	586	28,345
100x	673,335	23.81	1,146	28,279
200x	1,346,670	47.74	2,271	28,209